

# Assignment No.3

Title: Create a multiclass classification model to predict wine quality based on chemical properties.

Dataset: Wine Quality Dataset (UCI)

Name: Vaishnav Kalidas Temgire

Roll No: 23107127

Class : TY-B

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn as sk
```

```
In [3]: df = pd.read_csv("C:/Users/prath/Downloads/WineQT.csv")
```

```
In [4]: df
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alc
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
...	...	...	...	...	...	...	...	...	...	...	...
1138	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	
1139	6.8	0.620	0.08	1.9	0.068	28.0	38.0	0.99651	3.42	0.82	
1140	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	
1141	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	
1142	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	

1143 rows × 13 columns

```
In [5]: df.isnull().sum()
```

```
Out[5]: fixed acidity      0  
         volatile acidity   0  
         citric acid       0  
         residual sugar    0  
         chlorides          0  
         free sulfur dioxide 0  
         total sulfur dioxide 0  
         density            0  
         pH                 0  
         sulphates          0  
         alcohol             0  
         quality             0  
         Id                  0  
         dtype: int64
```

```
In [6]: df.isna().sum()
```

```
Out[6]: fixed acidity      0  
         volatile acidity   0  
         citric acid       0  
         residual sugar    0  
         chlorides          0  
         free sulfur dioxide 0  
         total sulfur dioxide 0  
         density            0  
         pH                 0  
         sulphates          0  
         alcohol             0  
         quality             0  
         Id                  0  
         dtype: int64
```

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1143 entries, 0 to 1142  
Data columns (total 13 columns):  
 #   Column           Non-Null Count Dtype  
 ---  ---  
 0   fixed acidity    1143 non-null  float64  
 1   volatile acidity 1143 non-null  float64  
 2   citric acid     1143 non-null  float64  
 3   residual sugar  1143 non-null  float64  
 4   chlorides        1143 non-null  float64  
 5   free sulfur dioxide 1143 non-null  float64  
 6   total sulfur dioxide 1143 non-null  float64  
 7   density          1143 non-null  float64  
 8   pH               1143 non-null  float64  
 9   sulphates        1143 non-null  float64  
 10  alcohol          1143 non-null  float64  
 11  quality          1143 non-null  int64  
 12  Id               1143 non-null  int64  
dtypes: float64(11), int64(2)  
memory usage: 116.2 KB
```


```
In [8]: df.describe()
```

	fixed acidity	volatile acidity	ci		
0	7.4	0.700	0.00	1.9	0.0
76	7.8	0.880	0.00	2.6	0.0
98	7.8	0.760	0.04	2.3	0.0
92	11.2	0.280	0.56	1.9	0.0
75	7.4	0.700	0.00	1.9	0.0
76	---	---	---	---	---
---	---	---	---	---	---
1138	6.3	0.510	0.13	2.3	0.0
76	6.8	0.620	0.08	1.9	0.0
68	6.2	0.600	0.08	2.0	0.0
1140	5.9	0.550	0.10	2.2	0.0
90	5.9	0.645	0.12	2.0	0.0
1141	5.9	0.550	0.10	2.2	0.0
62	5.9	0.645	0.12	2.0	0.0
1142	5.9	0.645	0.12	2.0	0.0
75	5.9	0.645	0.12	2.0	0.0
	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
---	---	---	---	---	---
1138	29.0	40.0	0.99574	3.42	0.75
1139	28.0	38.0	0.99651	3.42	0.82
1140	32.0	44.0	0.99490	3.45	0.58
1141	39.0	51.0	0.99512	3.52	0.76
1142	32.0	44.0	0.99547	3.57	0.71
	alcohol	quality	Id		
0	9.4	5	0		
1	9.8	5	1		
2	9.8	5	2		
3	9.8	6	3		
4	9.4	5	4		
---	---	---	---		
1138	11.0	6	1592		
1139	9.5	6	1593		
1140	10.5	5	1594		
1141	11.2	6	1595		
1142	10.2	5	1597		

[1143 rows x 13 columns]&gt;


In [9]: `sns.pairplot(df)`

Out[9]: &lt;seaborn.axisgrid.PairGrid at 0x1b3d44435d0&gt;




```
In [100]: df.hist(figsize=(12,10), bins=20, edgecolor="black")
plt.suptitle("Feature Distributions", fontsize=16)
plt.show()
```

## Feature Distributions



```
In [106]: plt.figure(figsize=(10,8))
sns.heatmap(df.corr(), annot=True)
plt.title("Correlation Heatmap")
plt.show()
```



```
In [10]: ## Dependent and Independent Variables
x = df.drop(columns=['quality','Id'],axis=1)
y= df['quality']
```

```
In [ ]: ## SMOTE For the Imbalance dataset

from imblearn.over_sampling import SMOTE
oversample = SMOTE()
x,y = oversample.fit_resample(x,df['quality'])
```

```
In [103...]: y.value_counts()
```

```
Out[103]: quality
5    483
6    483
7    483
4    483
8    483
3    483
Name: count, dtype: int64
```

```
In [105...]: x.shape,y.shape
```

```
Out[105]: ((2898, 11), (2898,))
```

```
In [60]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,train_size=0.8,random_
```

```
In [61]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
```

```
In [62]: x_train_scaled = sc.fit_transform(x_train)
```

```
In [63]: x_test_scaled = sc.transform(x_test)
```

```
In [119... from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
```

```
In [120... model
```

Out[120]:

```
In [121... model.fit(x_train_scaled,y_train)
```

Out[121]:

```
In [122... y_pred = model.predict(x_test_scaled)
```

```
In [123... y_pred
```

Out[123]:

```
array([8, 3, 5, 4, 7, 3, 7, 8, 5, 4, 6, 4, 3, 6, 3, 4, 6, 8, 8, 8, 7, 5,
       3, 6, 3, 8, 4, 5, 4, 3, 7, 4, 3, 5, 7, 7, 3, 4, 4, 4, 4, 3, 3, 3, 3,
       8, 4, 8, 5, 7, 8, 4, 3, 8, 6, 4, 4, 3, 3, 3, 4, 3, 5, 4, 6, 6, 6, 3,
       6, 7, 4, 5, 3, 5, 5, 3, 7, 6, 8, 5, 4, 5, 3, 8, 3, 4, 3, 8, 5, 4,
       7, 3, 7, 5, 5, 7, 3, 7, 4, 3, 7, 7, 6, 4, 3, 3, 4, 7, 7, 8, 4, 7,
       7, 5, 4, 6, 6, 7, 5, 7, 8, 4, 7, 4, 7, 6, 4, 4, 4, 4, 3, 7, 4, 4, 8,
       4, 3, 8, 7, 6, 4, 8, 3, 8, 8, 4, 6, 4, 3, 6, 4, 8, 5, 4, 7, 6, 7,
       4, 7, 8, 4, 4, 6, 3, 7, 4, 4, 8, 8, 8, 5, 5, 8, 3, 5, 3, 3, 5, 4,
       8, 5, 7, 7, 8, 5, 3, 7, 5, 6, 7, 7, 8, 3, 4, 6, 3, 6, 6, 6, 8, 6,
       6, 6, 5, 7, 6, 4, 5, 3, 3, 5, 7, 6, 8, 8, 4, 8, 5, 5, 7, 4, 4, 5,
       3, 8, 8, 5, 3, 7, 8, 7, 3, 3, 3, 8, 8, 4, 3, 7, 6, 6, 7, 4, 8, 5,
       4, 6, 7, 7, 5, 8, 7, 4, 7, 5, 4, 6, 4, 6, 7, 5, 4, 5, 4, 4, 6, 6,
       8, 5, 6, 3, 7, 7, 3, 6, 3, 3, 6, 5, 7, 4, 6, 4, 7, 6, 7, 6, 8, 8,
       4, 4, 5, 5, 4, 4, 6, 5, 5, 4, 6, 3, 5, 8, 5, 5, 8, 6, 6, 8, 3, 4,
       8, 8, 7, 6, 7, 3, 6, 7, 3, 6, 5, 8, 3, 7, 6, 6, 3, 3, 7, 3, 8, 7, 6,
       3, 5, 5, 8, 4, 8, 5, 3, 3, 6, 3, 4, 5, 3, 3, 5, 5, 4, 5, 8, 6,
       7, 3, 4, 3, 7, 3, 5, 7, 6, 8, 7, 5, 5, 4, 6, 3, 7, 6, 4, 7, 8, 6,
       4, 5, 4, 5, 8, 7, 7, 8, 7, 6, 7, 4, 8, 8, 4, 4, 5, 5, 3, 5,
       5, 7, 4, 6, 7, 3, 8, 5, 5, 6, 6, 5, 6, 3, 4, 3, 6, 5, 4, 3, 7, 5,
       4, 7, 7, 4, 7, 6, 8, 7, 5, 7, 8, 8, 3, 6, 5, 8, 7, 8, 5, 7, 5, 4,
       8, 8, 5, 6, 6, 3, 8, 4, 5, 5, 7, 8, 5, 8, 5, 6, 3, 7, 6, 6, 7, 8,
       6, 8, 8, 8, 8, 7, 3, 5, 8, 6, 4, 5, 5, 6, 6, 4, 7, 7, 5, 6, 8, 6,
       6, 5, 4, 6, 8, 6, 8, 7, 8, 7, 3, 3, 4, 6, 8, 5, 5, 6, 4, 7, 6, 8,
       4, 8, 3, 3, 6, 3, 5, 6, 7, 8, 4, 3, 3, 6, 8, 6, 8, 5, 4, 5, 5, 4,
       4, 6, 8, 5, 8, 7, 8, 7, 6, 7, 8, 6, 8, 7, 3, 6, 8, 7, 5, 5, 8, 7,
       6, 6, 3, 7, 8, 3, 7, 5, 6, 7, 7, 7, 8, 5, 4, 7, 5, 8, 3, 4, 8, 3,
       4, 7, 3, 6, 6, 7, 3, 4, 4], dtype=int64)
```

```
In [124... from sklearn.metrics import classification_report,accuracy_score,precision_
cr = classification_report(y_test,y_pred)
acc_dt = accuracy_score(y_test,y_pred)
re = recall_score(y_test,y_pred,average='micro')
```

```
f1 = f1_score(y_test,y_pred,average='micro')
pr =precision_score(y_test,y_pred,average='micro')
```

In [125]: `print(cr)`

	precision	recall	f1-score	support
3	0.95	0.91	0.93	97
4	0.79	0.84	0.81	92
5	0.68	0.59	0.63	108
6	0.54	0.54	0.54	95
7	0.79	0.84	0.81	96
8	0.91	0.96	0.93	92
accuracy			0.77	580
macro avg	0.77	0.78	0.78	580
weighted avg	0.77	0.77	0.77	580

In [126]: `re`

Out[126]: 0.7741379310344828

In [127]: `pr`

Out[127]: 0.7741379310344828

In [128]: `acc_dt`

Out[128]: 0.7741379310344828


In [129]: `f1`

Out[129]: 0.7741379310344828

In [130]: `## Random Forest Classifier`

```
from sklearn.ensemble import RandomForestClassifier
rc = RandomForestClassifier()
```

In [131]: `rc`

Out[131]: 

► Parameters

In [132]: `rc.fit(x_train_scaled,y_train)`

Out[132]: 

► Parameters

In [133]: `y_pred = rc.predict(x_test_scaled)`

In [134]: `y_pred`

```
Out[134]: array([7, 3, 5, 4, 7, 3, 7, 8, 5, 4, 6, 7, 3, 5, 3, 6, 6, 8, 8, 8, 8, 7, 5,
   3, 6, 3, 8, 5, 3, 4, 3, 7, 4, 3, 3, 6, 7, 3, 4, 4, 4, 4, 3, 3, 3, 4,
   8, 4, 8, 6, 7, 8, 4, 3, 7, 6, 4, 4, 3, 3, 3, 4, 3, 5, 4, 6, 5, 3,
   5, 7, 4, 7, 3, 5, 4, 3, 7, 6, 8, 6, 4, 4, 3, 8, 3, 4, 3, 8, 4, 4,
   7, 3, 7, 6, 5, 7, 3, 7, 4, 3, 6, 7, 6, 4, 3, 3, 4, 7, 7, 8, 5, 7,
   6, 5, 6, 6, 5, 8, 3, 7, 8, 4, 7, 4, 7, 4, 4, 6, 4, 3, 7, 4, 4, 8,
   6, 4, 8, 7, 5, 4, 8, 3, 8, 8, 4, 4, 4, 4, 3, 6, 4, 8, 4, 4, 7, 6, 7,
   4, 7, 8, 4, 4, 5, 3, 7, 6, 5, 8, 8, 8, 5, 6, 7, 3, 5, 3, 3, 5, 6,
   6, 5, 7, 7, 8, 6, 3, 7, 5, 6, 7, 4, 8, 3, 6, 5, 3, 6, 6, 6, 8, 4,
   6, 6, 5, 7, 5, 4, 3, 3, 3, 5, 7, 6, 8, 8, 5, 8, 5, 4, 7, 5, 5, 5,
   3, 8, 8, 5, 3, 7, 8, 7, 4, 3, 3, 8, 8, 4, 3, 7, 7, 6, 7, 4, 8, 5,
   4, 6, 7, 7, 3, 8, 7, 4, 7, 5, 4, 5, 4, 5, 7, 5, 4, 5, 3, 5, 6, 6,
   8, 5, 6, 3, 6, 7, 3, 5, 3, 3, 6, 6, 7, 4, 7, 4, 8, 7, 7, 6, 7, 8,
   5, 4, 6, 6, 5, 4, 6, 6, 5, 4, 6, 3, 5, 8, 4, 6, 8, 5, 6, 8, 3, 4,
   8, 8, 7, 6, 7, 3, 6, 7, 3, 5, 5, 6, 3, 7, 6, 5, 3, 7, 3, 8, 7, 5,
   3, 5, 5, 8, 4, 8, 5, 3, 3, 6, 3, 4, 4, 3, 3, 5, 6, 4, 5, 8, 5,
   7, 3, 6, 3, 7, 3, 5, 7, 5, 8, 7, 5, 5, 4, 7, 3, 7, 6, 4, 7, 8, 6,
   4, 5, 4, 5, 6, 8, 7, 8, 7, 7, 6, 4, 8, 8, 5, 4, 3, 3, 3, 5,
   6, 7, 4, 7, 7, 3, 8, 5, 4, 6, 6, 6, 5, 3, 4, 3, 6, 5, 4, 4, 7, 5,
   4, 7, 7, 4, 7, 6, 7, 7, 5, 7, 8, 8, 3, 6, 5, 8, 7, 8, 5, 7, 4, 4,
   8, 8, 5, 5, 6, 3, 8, 4, 5, 5, 7, 8, 5, 8, 5, 6, 3, 6, 6, 7, 7, 8,
   6, 8, 8, 8, 8, 7, 3, 5, 8, 6, 4, 5, 5, 6, 7, 4, 8, 7, 6, 6, 8, 6,
   5, 7, 4, 7, 8, 5, 8, 7, 8, 7, 3, 3, 4, 5, 8, 5, 6, 7, 4, 7, 7, 8,
   4, 8, 3, 3, 5, 4, 5, 5, 7, 8, 4, 3, 3, 6, 8, 5, 8, 5, 4, 3, 4, 4,
   4, 6, 8, 5, 8, 7, 8, 7, 6, 6, 8, 4, 8, 6, 3, 7, 8, 7, 5, 6, 8, 7,
   6, 6, 3, 7, 8, 3, 7, 5, 4, 7, 7, 7, 8, 5, 4, 7, 5, 8, 3, 4, 8, 3,
   4, 7, 3, 5, 6, 7, 3, 5], dtype=int64)
```

```
In [135...]: from sklearn.metrics import classification_report,accuracy_score,precision_
cr = classification_report(y_test,y_pred)
acc_rf = accuracy_score(y_test,y_pred)
re = recall_score(y_test,y_pred,average='micro')
f1 = f1_score(y_test,y_pred,average='micro')
pr = precision_score(y_test,y_pred,average='micro')
cm = confusion_matrix(y_test,y_pred)
```

```
In [136...]: print(cr)
```

	precision	recall	f1-score	support
3	1.00	1.00	1.00	97
4	0.91	0.97	0.94	92
5	0.80	0.70	0.75	108
6	0.62	0.58	0.60	95
7	0.81	0.92	0.86	96
8	0.99	1.00	0.99	92
accuracy			0.86	580
macro avg	0.85	0.86	0.86	580
weighted avg	0.85	0.86	0.85	580

```
In [137...]: acc_rf
```

```
Out[137]: 0.8568965517241379
```

```
In [138...]: re
```

```
Out[138]: 0.8568965517241379
```

```
In [139...]: f1
```


In [139]: pr  
Out[139]: 0.8568965517241379

In [140... pr  
Out[140]: 0.8568965517241379

In [141... print(cm)


```
[[97  0  0  0  0  0]
 [ 0 89  1  2  0  0]
 [ 0  3 76 24  5  0]
 [ 0  6 18 55 16  0]
 [ 0  0  0  7 88  1]
 [ 0  0  0  0  0 92]]
```

In [142... from sklearn.metrics import ConfusionMatrixDisplay  
disp = ConfusionMatrixDisplay(confusion\_matrix=cm, display\_labels=model.classes\_)  
disp.plot()  
plt.title("Confusion Matrix")  
plt.show()




In [143... from sklearn.metrics import roc\_curve, auc  
from sklearn.preprocessing import label\_binarize  
  
y\_test\_bin = label\_binarize(y\_test, classes=sorted(df['quality'].unique()))  
  
y\_score = rc.predict\_proba(x\_test\_scaled)  
  
plt.figure(figsize=(6,5))  
for i in range(y\_test\_bin.shape[1]):  
 fpr, tpr, \_ = roc\_curve(y\_test\_bin[:, i], y\_score[:, i])  
 plt.plot(fpr, tpr, label=f"Class {i}")  
  
plt.plot([0,1], [0,1], "k--")

```
plt.xlabel("FPR")
plt.ylabel("TPR")
plt.title("ROC - Random Forest")
plt.legend()
plt.show()
```



```
In [145]: plt.figure(figsize=(6,5))
plt.bar(["Decision Tree", "Random Forest"], [acc_dt, acc_rf], color=["blue", "red"])
plt.ylabel("Accuracy")
plt.title("Model Accuracy Comparison")
plt.show()
```

### Model Accuracy Comparison



In [ ]: