

Assignment No-2

Title-Develop a Ridge and Lasso regression model to predict the number of bike rentals based on weather conditions and time. Dataset: Bike Sharing Dataset (UCI)

Name: Vaishnav Kalidas Temgire

Roll No:23107127 **Class** : TY-B **Batch** : B

```
In [66]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [67]: df = pd.read_csv("hour.csv")
df
```

	instant	dteday	season	yr	mnth	hr	holiday	weekday	workingday
0	1	2011-01-01	1	0	1	0	0		6
1	2	2011-01-01	1	0	1	1	0		6
2	3	2011-01-01	1	0	1	2	0		6
3	4	2011-01-01	1	0	1	3	0		6
4	5	2011-01-01	1	0	1	4	0		6
...
17374	17375	2012-12-31	1	1	12	19	0		1
17375	17376	2012-12-31	1	1	12	20	0		1
17376	17377	2012-12-31	1	1	12	21	0		1
17377	17378	2012-12-31	1	1	12	22	0		1
17378	17379	2012-12-31	1	1	12	23	0		1

17379 rows × 17 columns

```
In [68]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 17379 entries, 0 to 17378
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   instant     17379 non-null   int64  
 1   dteday      17379 non-null   object  
 2   season      17379 non-null   int64  
 3   yr          17379 non-null   int64  
 4   mnth        17379 non-null   int64  
 5   hr          17379 non-null   int64  
 6   holiday     17379 non-null   int64  
 7   weekday     17379 non-null   int64  
 8   workingday  17379 non-null   int64  
 9   weathersit  17379 non-null   int64  
 10  temp         17379 non-null   float64 
 11  atemp        17379 non-null   float64 
 12  hum          17379 non-null   float64 
 13  windspeed    17379 non-null   float64 
 14  casual       17379 non-null   int64  
 15  registered   17379 non-null   int64  
 16  cnt          17379 non-null   int64  
dtypes: float64(4), int64(12), object(1)
memory usage: 2.3+ MB
```

```
In [69]: df.size
```

```
Out[69]: 295443
```

```
In [70]: df.shape
```

```
Out[70]: (17379, 17)
```

```
In [71]: df.ndim
```

```
Out[71]: 2
```

```
In [72]: df.describe()
```

Out[72]:

	instant	season	yr	mnth	hr	
count	17379.0000	17379.000000	17379.000000	17379.000000	17379.000000	1737
mean	8690.0000	2.501640	0.502561	6.537775	11.546752	
std	5017.0295	1.106918	0.500008	3.438776	6.914405	
min	1.0000	1.000000	0.000000	1.000000	0.000000	
25%	4345.5000	2.000000	0.000000	4.000000	6.000000	
50%	8690.0000	3.000000	1.000000	7.000000	12.000000	
75%	13034.5000	3.000000	1.000000	10.000000	18.000000	
max	17379.0000	4.000000	1.000000	12.000000	23.000000	

In [73]: `df.columns`

Out[73]: `Index(['instant', 'dteday', 'season', 'yr', 'mnth', 'hr', 'holiday', 'weekday', 'workingday', 'weathersit', 'temp', 'atemp', 'hum', 'windspeed', 'casual', 'registered', 'cnt'], dtype='object')`

In [74]: `df.isna().sum()`

```
Out[74]: 0
instant 0
dteday 0
season 0
yr 0
mnth 0
hr 0
holiday 0
weekday 0
workingday 0
weathersit 0
temp 0
atemp 0
hum 0
windspeed 0
casual 0
registered 0
cnt 0
```

dtype: int64

```
In [75]: df.notnull().sum()
```

```
Out[75]: 0
instant 17379
dteday 17379
season 17379
yr 17379
mnth 17379
hr 17379
holiday 17379
weekday 17379
workingday 17379
weathersit 17379
temp 17379
atemp 17379
hum 17379
windspeed 17379
casual 17379
registered 17379
cnt 17379
```

dtype: int64

```
In [76]: df.min()
```

```
Out[76]:
```

	0
instant	1
dteday	2011-01-01
season	1
yr	0
mnth	1
hr	0
holiday	0
weekday	0
workingday	0
weathersit	1
temp	0.02
atemp	0.0
hum	0.0
windspeed	0.0
casual	0
registered	0
cnt	1

dtype: object

```
In [77]: df["cnt"].max()
```

```
Out[77]: 977
```

```
In [78]: df['registered'].var()
```

```
Out[78]: 22909.027998823447
```

```
In [79]: df['cnt'].mean()
```

```
Out[79]: np.float64(189.46308763450142)
```

```
In [80]: df['cnt'].median()
```

```
Out[80]: 142.0
```

```
In [81]: df['cnt'].mode()
```

Out[81]:

	cnt
0	5

dtype: int64

In [82]:

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['dteday_encoded'] = le.fit_transform(df['dteday'])
print(df[['dteday', 'dteday_encoded']].head())
```

	dteday	dteday_encoded
0	2011-01-01	0
1	2011-01-01	0
2	2011-01-01	0
3	2011-01-01	0
4	2011-01-01	0

In [83]:

Out[83]:

	instant	dteday	season	yr	mnth	hr	holiday	weekday	workingday
0	1	2011-01-01	1	0	1	0	0	6	
1	2	2011-01-01	1	0	1	1	0	6	
2	3	2011-01-01	1	0	1	2	0	6	
3	4	2011-01-01	1	0	1	3	0	6	
4	5	2011-01-01	1	0	1	4	0	6	
...
17374	17375	2012-12-31	1	1	12	19	0	1	
17375	17376	2012-12-31	1	1	12	20	0	1	
17376	17377	2012-12-31	1	1	12	21	0	1	
17377	17378	2012-12-31	1	1	12	22	0	1	
17378	17379	2012-12-31	1	1	12	23	0	1	

17379 rows × 18 columns

In [84]:

```
x = df.drop(['cnt','dteday'], axis=1)
print(x)
```

```

      instant  season  yr  mnth  hr  holiday  weekday  workingday  \
0            1       1  0     1   0        0         6           0
1            2       1  0     1   1        0         6           0
2            3       1  0     1   2        0         6           0
3            4       1  0     1   3        0         6           0
4            5       1  0     1   4        0         6           0
...
17374      17375    1   1   12   19        0         1           1
17375      17376    1   1   12   20        0         1           1
17376      17377    1   1   12   21        0         1           1
17377      17378    1   1   12   22        0         1           1
17378      17379    1   1   12   23        0         1           1

      weathersit  temp  atemp  hum  windspeed  casual  registered  \
0            1  0.24  0.2879  0.81  0.0000      3        13
1            1  0.22  0.2727  0.80  0.0000      8        32
2            1  0.22  0.2727  0.80  0.0000      5        27
3            1  0.24  0.2879  0.75  0.0000      3        10
4            1  0.24  0.2879  0.75  0.0000      0         1
...
17374      2  0.26  0.2576  0.60  0.1642     11       108
17375      2  0.26  0.2576  0.60  0.1642      8        81
17376      1  0.26  0.2576  0.60  0.1642      7        83
17377      1  0.26  0.2727  0.56  0.1343     13       48
17378      1  0.26  0.2727  0.65  0.1343     12       37

      dteday_encoded
0                  0
1                  0
2                  0
3                  0
4                  0
...
17374            730
17375            730
17376            730
17377            730
17378            730

```

[17379 rows x 16 columns]

In [85]: `y = df['cnt']
print(y)`

```
0      16
1      40
2      32
3      13
4      1
...
17374    119
17375     89
17376     90
17377     61
17378     49
Name: cnt, Length: 17379, dtype: int64
```

```
In [86]: # 1) Linear Regression Model
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, rando
```

```
In [87]: # Scaling --
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()

x_train_scaled = sc.fit_transform(x_train)
x_test_scaled = sc.fit_transform(x_test)
```

```
In [88]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(x_train_scaled, y_train)
```

```
Out[88]: ▾ LinearRegression ⓘ ⓘ
LinearRegression()
```

```
In [89]: y_pred = model.predict(x_test_scaled)
y_pred
```

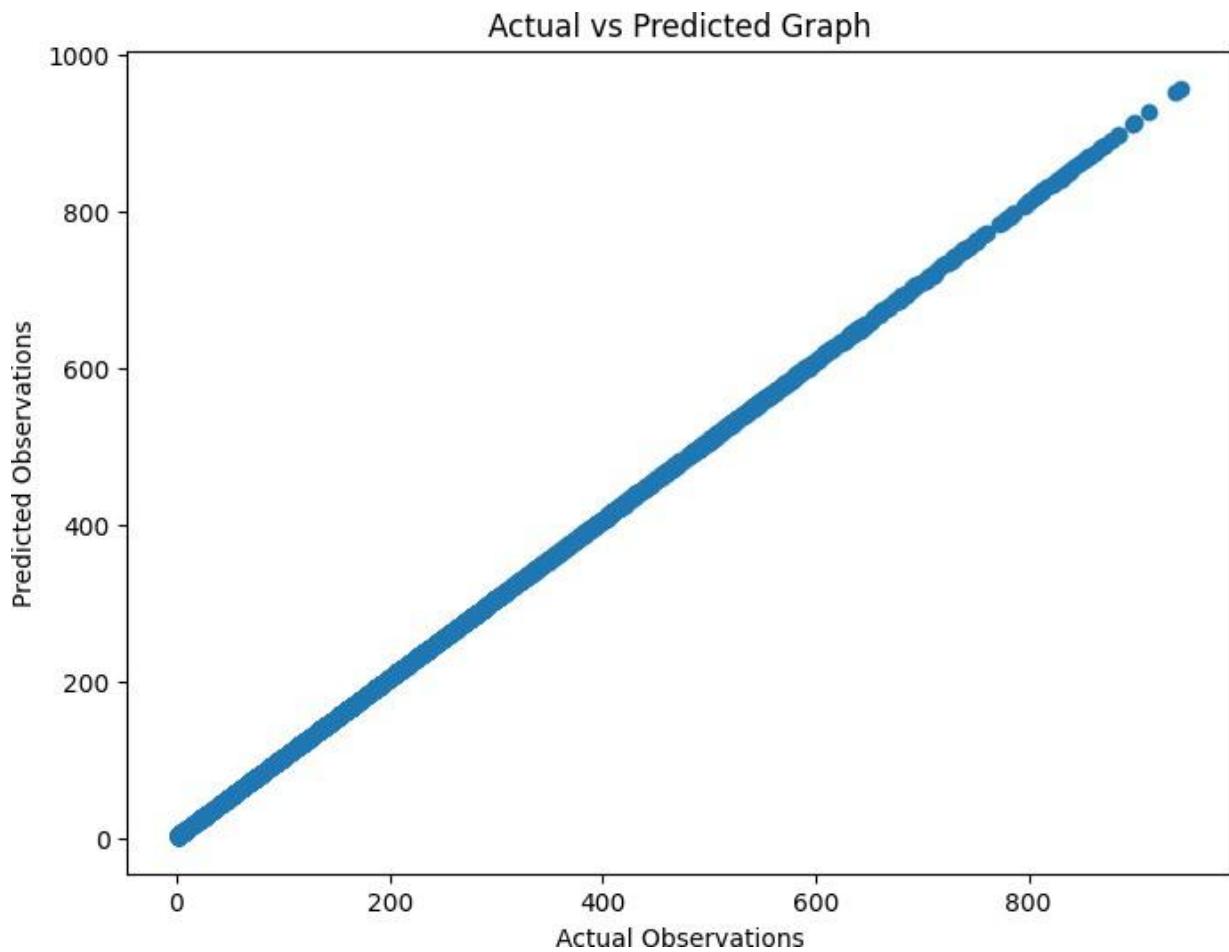
```
Out[89]: array([270.66269285, 40.33412021, 225.10554443, ..., 8.87444699,
   101.85933915, 116.10363989])
```

```
In [90]: from sklearn.metrics import r2_score, mean_absolute_error
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
print("R^2 : ",r2)
print("Mean Absolute Error : ",mae)
```

```
R^2 : 0.9995354285453604
Mean Absolute Error : 3.1236835097766202
```

```
In [91]: plt.figure(figsize=(8,6))
plt.scatter(y_test,y_pred)
plt.title("Actual vs Predicted Graph")
plt.xlabel("Actual Observations")
plt.ylabel("Predicted Observations")
```

```
plt.show()
```



In [92]: # Ridge

In [93]: `from sklearn.linear_model import Ridge`
R_model = Ridge()

In [94]: R_model.fit(x_train_scaled,y_train)

Out[94]:
▼ Ridge ⓘ ⓘ
Ridge()

In [95]: y_pred_R = R_model.predict(x_test_scaled)
y_pred_R

Out[95]: array([270.66179557, 40.34929279, 225.11821542, ..., 8.88925947,
101.8621773 , 116.10228523])

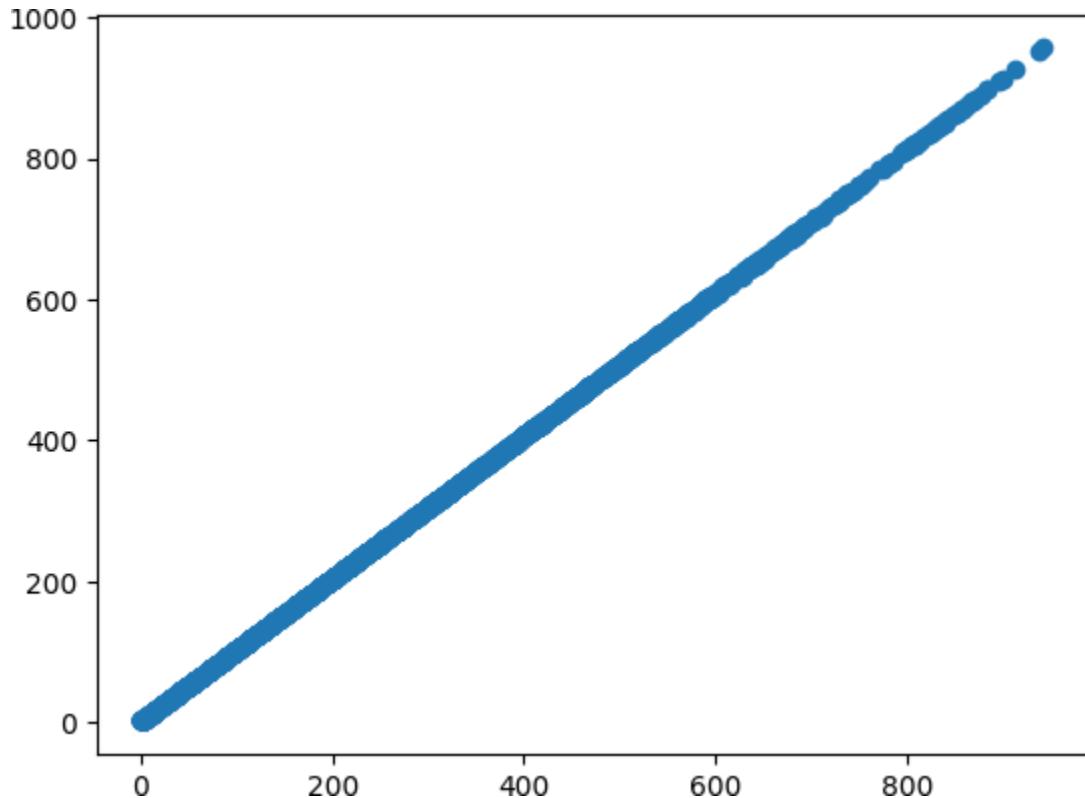
In [96]: `from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score`
mse = mean_squared_error(y_test, y_pred_R)
mae = mean_absolute_error(y_test, y_pred_R)

```
r2 = r2_score(y_test, y_pred_R)
print("mean_squared_error : ",mse)
print("mean_absolute_error : ",mae)
print("r2_score : ",r2)
```

```
mean_squared_error :  14.935141607095735
mean_absolute_error :  3.123683509776621
r2_score :  0.999537187797581
```

```
In [97]: plt.scatter(y_test,y_pred_R)
```

```
Out[97]: <matplotlib.collections.PathCollection at 0x7eebc0a50050>
```



```
In [98]: from sklearn.linear_model import Lasso
model_L = Lasso(alpha=1)
```

```
In [99]: model_L.fit(x_train_scaled,y_train)
```

```
Out[99]: ▾ Lasso ⓘ ⓘ
Lasso(alpha=1)
```

```
In [102...]: y_pred_L = model_L.predict(x_test_scaled)
y_pred_L
```

```
Out[102...]: array([269.83296506, 41.29329561, 224.5564963 , ..., 9.97068948,
   102.26735186, 116.47698844])
```

```
In [106]: from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score
mae = mean_absolute_error(y_test,y_pred_L)
mse = mean_squared_error(y_test,y_pred_L)
rmse = np.sqrt(mse)
print("mean absolute error : ",mae)
print("mean squared error : ",mse)
print("Root mean squared error : ",rmse)
r2 = r2_score(y_test, y_pred_L)
print("R^2 : ",r2)
```

mean absolute error : 11.634135792661272
 mean squared error : 3.1236835097766207
 Root mean squared error : 3.410884898770592
 R^2 : 0.9996394798153849

Day Dataset

```
In [112]: df1=pd.read_csv('day.csv')
df1
```

Out[112]:

	instant	dteday	season	yr	mnth	holiday	weekday	workingday	wea
0	1	2011-01-01	1	0	1	0	6		0
1	2	2011-01-02	1	0	1	0	0		0
2	3	2011-01-03	1	0	1	0	1		1
3	4	2011-01-04	1	0	1	0	2		1
4	5	2011-01-05	1	0	1	0	3		1
...
726	727	2012-12-27	1	1	12	0	4		1
727	728	2012-12-28	1	1	12	0	5		1
728	729	2012-12-29	1	1	12	0	6		0
729	730	2012-12-30	1	1	12	0	0		0
730	731	2012-12-31	1	1	12	0	1		1

731 rows × 16 columns

```
In [113]: df1.dtypes
```

```
Out[113...]
```

	0
instant	int64
dteday	object
season	int64
yr	int64
mnth	int64
holiday	int64
weekday	int64
workingday	int64
weathersit	int64
temp	float64
atemp	float64
hum	float64
windspeed	float64
casual	int64
registered	int64
cnt	int64

dtype: object

```
In [114...]: df1.isnull().sum()
```

```
Out[114...      0
     instant   0
     dteday    0
     season    0
     yr        0
     mnth     0
     holiday   0
     weekday   0
     workingday 0
     weathersit 0
     temp      0
     atemp     0
     hum       0
     windspeed 0
     casual    0
     registered 0
     cnt       0
```

dtype: int64

```
In [115... df1.ndim
```

```
Out[115... 2
```

```
In [116... df1.size
```

```
Out[116... 11696
```

```
In [117... df1.shape
```

```
Out[117... (731, 16)
```

```
In [118... df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731 entries, 0 to 730
Data columns (total 16 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   instant     731 non-null    int64  
 1   dteday      731 non-null    object  
 2   season      731 non-null    int64  
 3   yr          731 non-null    int64  
 4   mnth        731 non-null    int64  
 5   holiday     731 non-null    int64  
 6   weekday     731 non-null    int64  
 7   workingday  731 non-null    int64  
 8   weathersit  731 non-null    int64  
 9   temp         731 non-null    float64 
 10  atemp        731 non-null    float64 
 11  hum          731 non-null    float64 
 12  windspeed   731 non-null    float64 
 13  casual       731 non-null    int64  
 14  registered   731 non-null    int64  
 15  cnt          731 non-null    int64  
dtypes: float64(4), int64(11), object(1)
memory usage: 91.5+ KB
```

```
In [119... df1.describe()
```

```
Out[119...

|              | instant    | season     | yr         | mnth       | holiday    | weekday    |
|--------------|------------|------------|------------|------------|------------|------------|
| <b>count</b> | 731.000000 | 731.000000 | 731.000000 | 731.000000 | 731.000000 | 731.000000 |
| <b>mean</b>  | 366.000000 | 2.496580   | 0.500684   | 6.519836   | 0.028728   | 2.997264   |
| <b>std</b>   | 211.165812 | 1.110807   | 0.500342   | 3.451913   | 0.167155   | 2.004787   |
| <b>min</b>   | 1.000000   | 1.000000   | 0.000000   | 1.000000   | 0.000000   | 0.000000   |
| <b>25%</b>   | 183.500000 | 2.000000   | 0.000000   | 4.000000   | 0.000000   | 1.000000   |
| <b>50%</b>   | 366.000000 | 3.000000   | 1.000000   | 7.000000   | 0.000000   | 3.000000   |
| <b>75%</b>   | 548.500000 | 3.000000   | 1.000000   | 10.000000  | 0.000000   | 5.000000   |
| <b>max</b>   | 731.000000 | 4.000000   | 1.000000   | 12.000000  | 1.000000   | 6.000000   |


```

```
In [120... data=df.copy()
```

```
In [121... data=data.drop(['dteday'], axis=1)
data=data.drop(['instant'], axis=1)
data=data.drop(['atemp'], axis=1)
data=data.drop(['registered'], axis=1)
data=data.drop(['casual'], axis=1)
```

```
In [122... x=data.drop(['cnt'], axis=1)
y=data['cnt']
```

```
In [123...]: x_train, x_test, y_train, y_test = train_test_split(  
x, y, test_size=0.3, random_state=42)  
scaler=StandardScaler()  
x_train_scaled=scaler.fit_transform(x_train)  
x_test_scaled=scaler.fit_transform(x_test)
```

```
In [124...]: lr=LinearRegression()
```

```
In [127...]: lr.fit(x_train_scaled, y_train)  
y_pred=lr.predict(x_test_scaled)
```

```
In [135...]: mae = mean_absolute_error(y_test,y_pred)  
mse = mean_squared_error(y_test,y_pred)  
rmse = np.sqrt(mse)  
print("mean absolute error : ",mse)  
print("mean squared error : ",mae)  
print("Root mean squared error : ",rmse)  
r2 = r2_score(y_test, y_pred)  
print("R^2 : ",r2)
```

```
mean absolute error : 733994.9324822383  
mean squared error : 635.7374845148096  
Root mean squared error : 856.7350421701206  
R^2 : 0.8155312414571874
```

```
In [132...]: ridge=Ridge(alpha=10)
```

```
In [133...]: ridge.fit(x_train_scaled, y_train)
```

```
Out[133...]:  
▼ Ridge ⓘ ⓘ  
Ridge(alpha=10)
```

```
In [134...]: ridge_pred=ridge.predict(x_test_scaled)
```

```
In [136...]: mae = mean_absolute_error(y_test, ridge_pred)  
mse = mean_squared_error(y_test, ridge_pred)  
rmse = np.sqrt(mse)  
print("mean absolute error : ",mse)  
print("mean squared error : ",mae)  
print("Root mean squared error : ",rmse)  
r2 = r2_score(y_test, ridge_pred)  
print("R^2 : ",r2)
```

```
mean absolute error : 736023.3441440698  
mean squared error : 639.5834409528549  
Root mean squared error : 857.9180288023267  
R^2 : 0.8150214578544499
```

```
In [146...]: from sklearn.linear_model import Lasso  
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score  
from sklearn.preprocessing import StandardScaler
```

```
import numpy as np
```

```
In [147]: scaler = StandardScaler()
x_train_scaled = scaler.fit_transform(x_train)
x_test_scaled = scaler.transform(x_test)
```

```
In [148]: lasso = Lasso(alpha=1)
lasso.fit(x_train_scaled, y_train)
lasso_pred = lasso.predict(x_test_scaled)
```

```
In [149]: mae = mean_absolute_error(y_test, lasso_pred)
mse = mean_squared_error(y_test, lasso_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, lasso_pred)

print("Mean Absolute Error:", mae)
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R^2:", r2)
```

```
Mean Absolute Error: 631.8713036109639
Mean Squared Error: 703468.622956881
Root Mean Squared Error: 838.7303636788649
R^2: 0.8232031614825652
```

```
In [ ]:
```